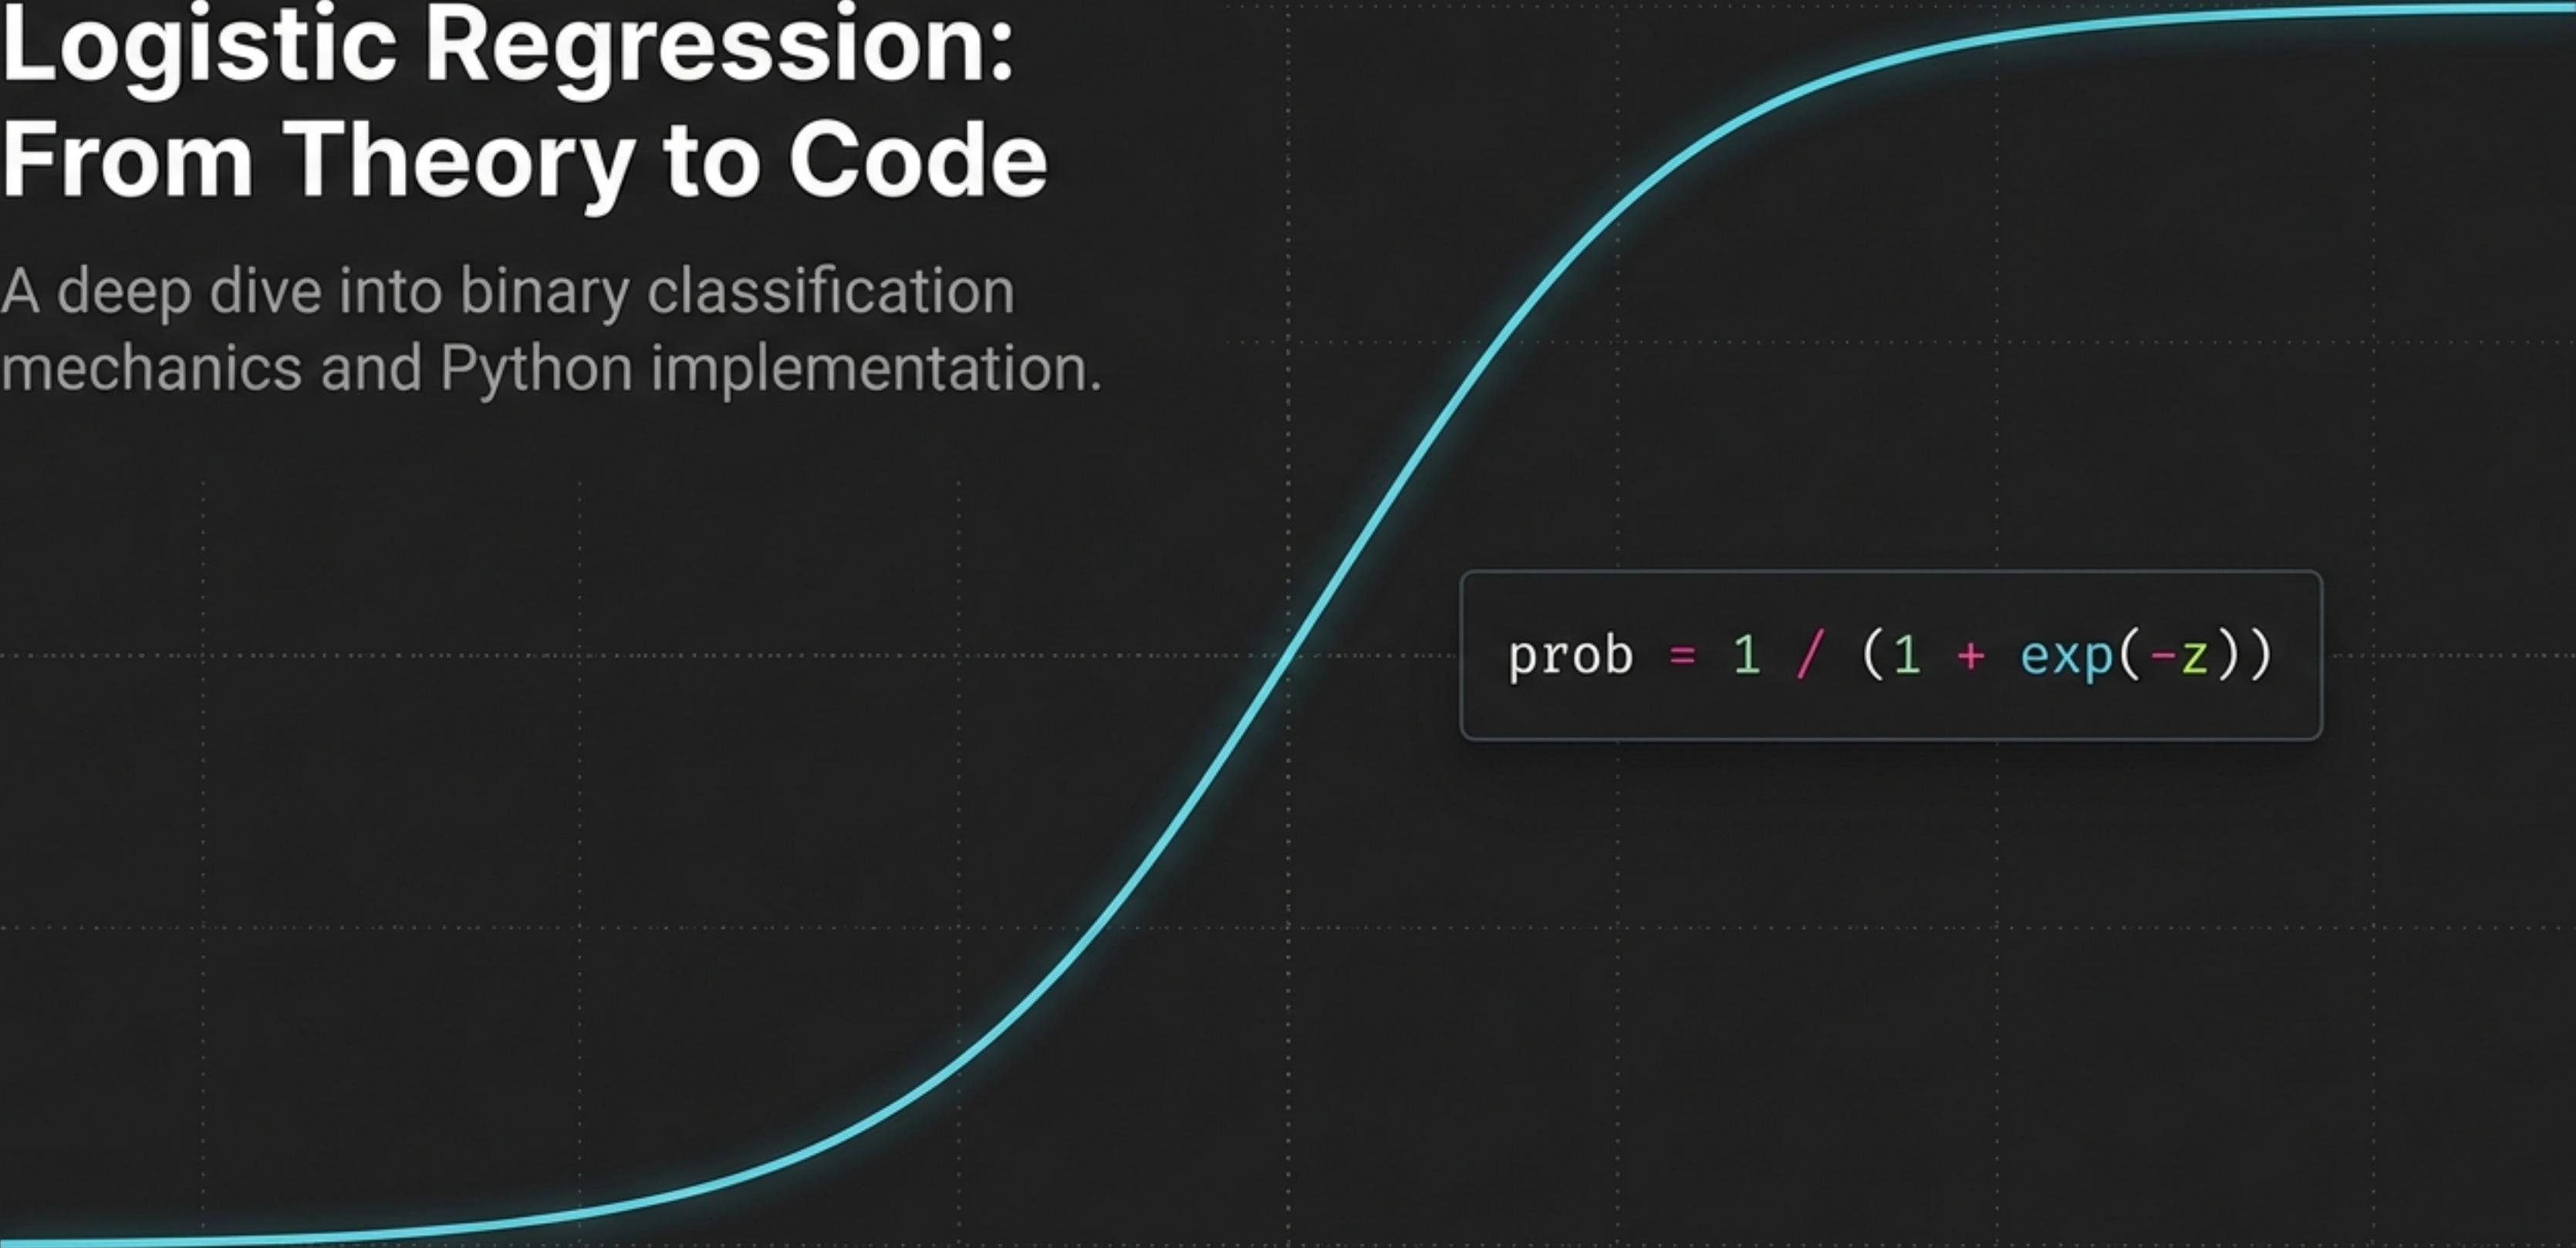


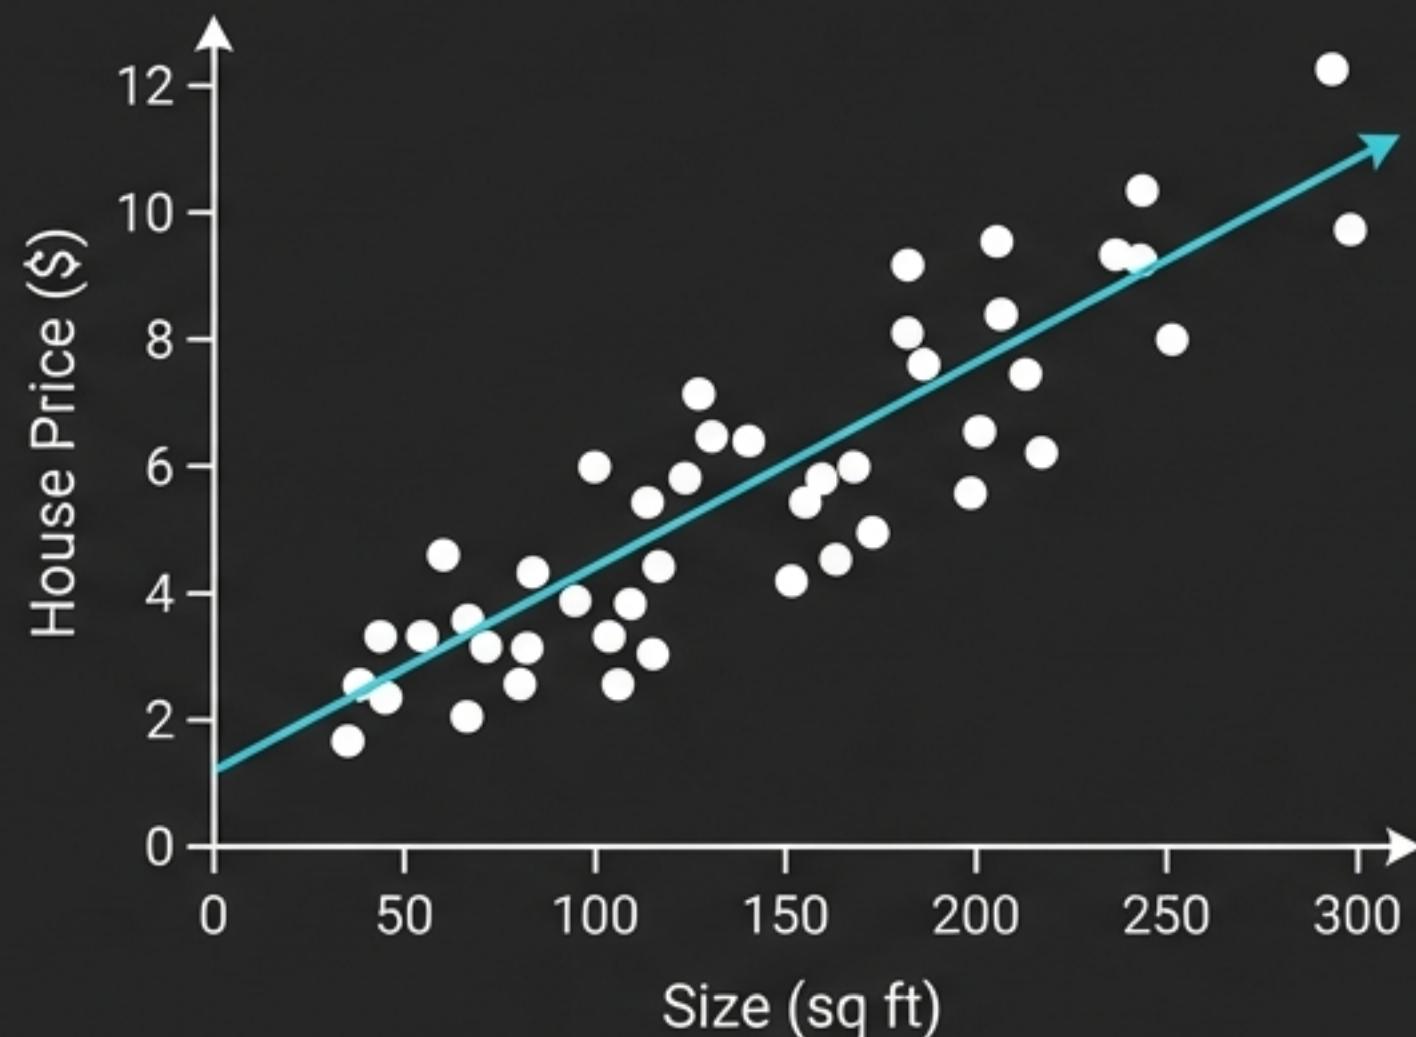
Logistic Regression: From Theory to Code

A deep dive into binary classification mechanics and Python implementation.

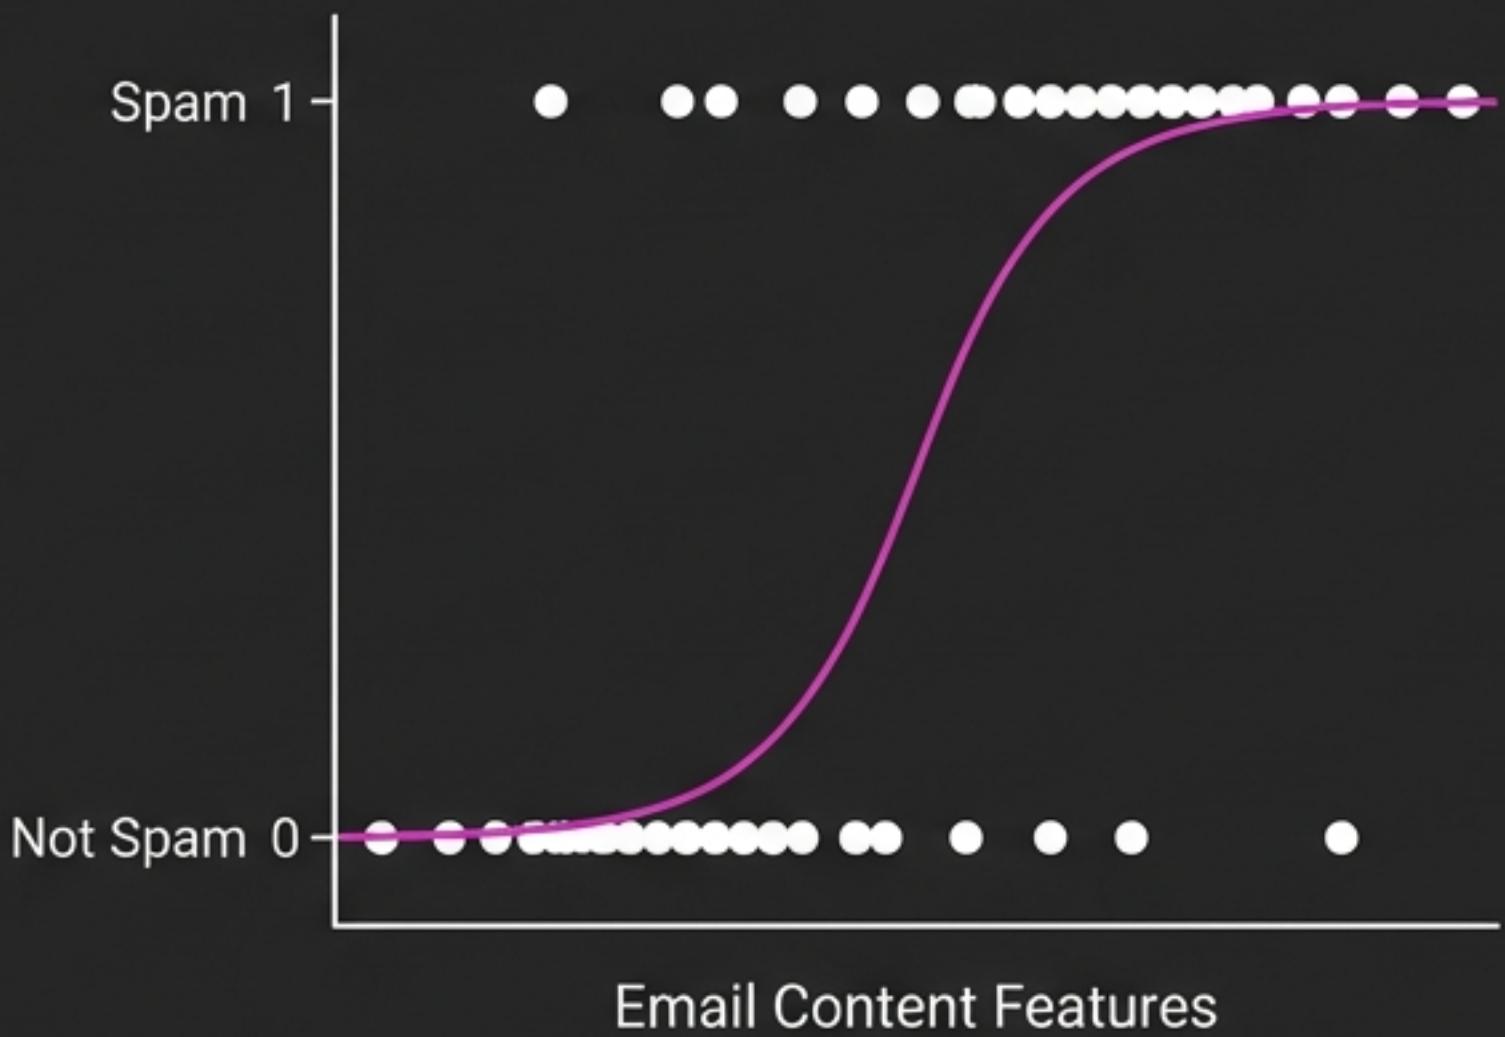

$$\text{prob} = 1 / (1 + \exp(-z))$$

The Shift from Regression to Classification

Linear Regression (Continuous)



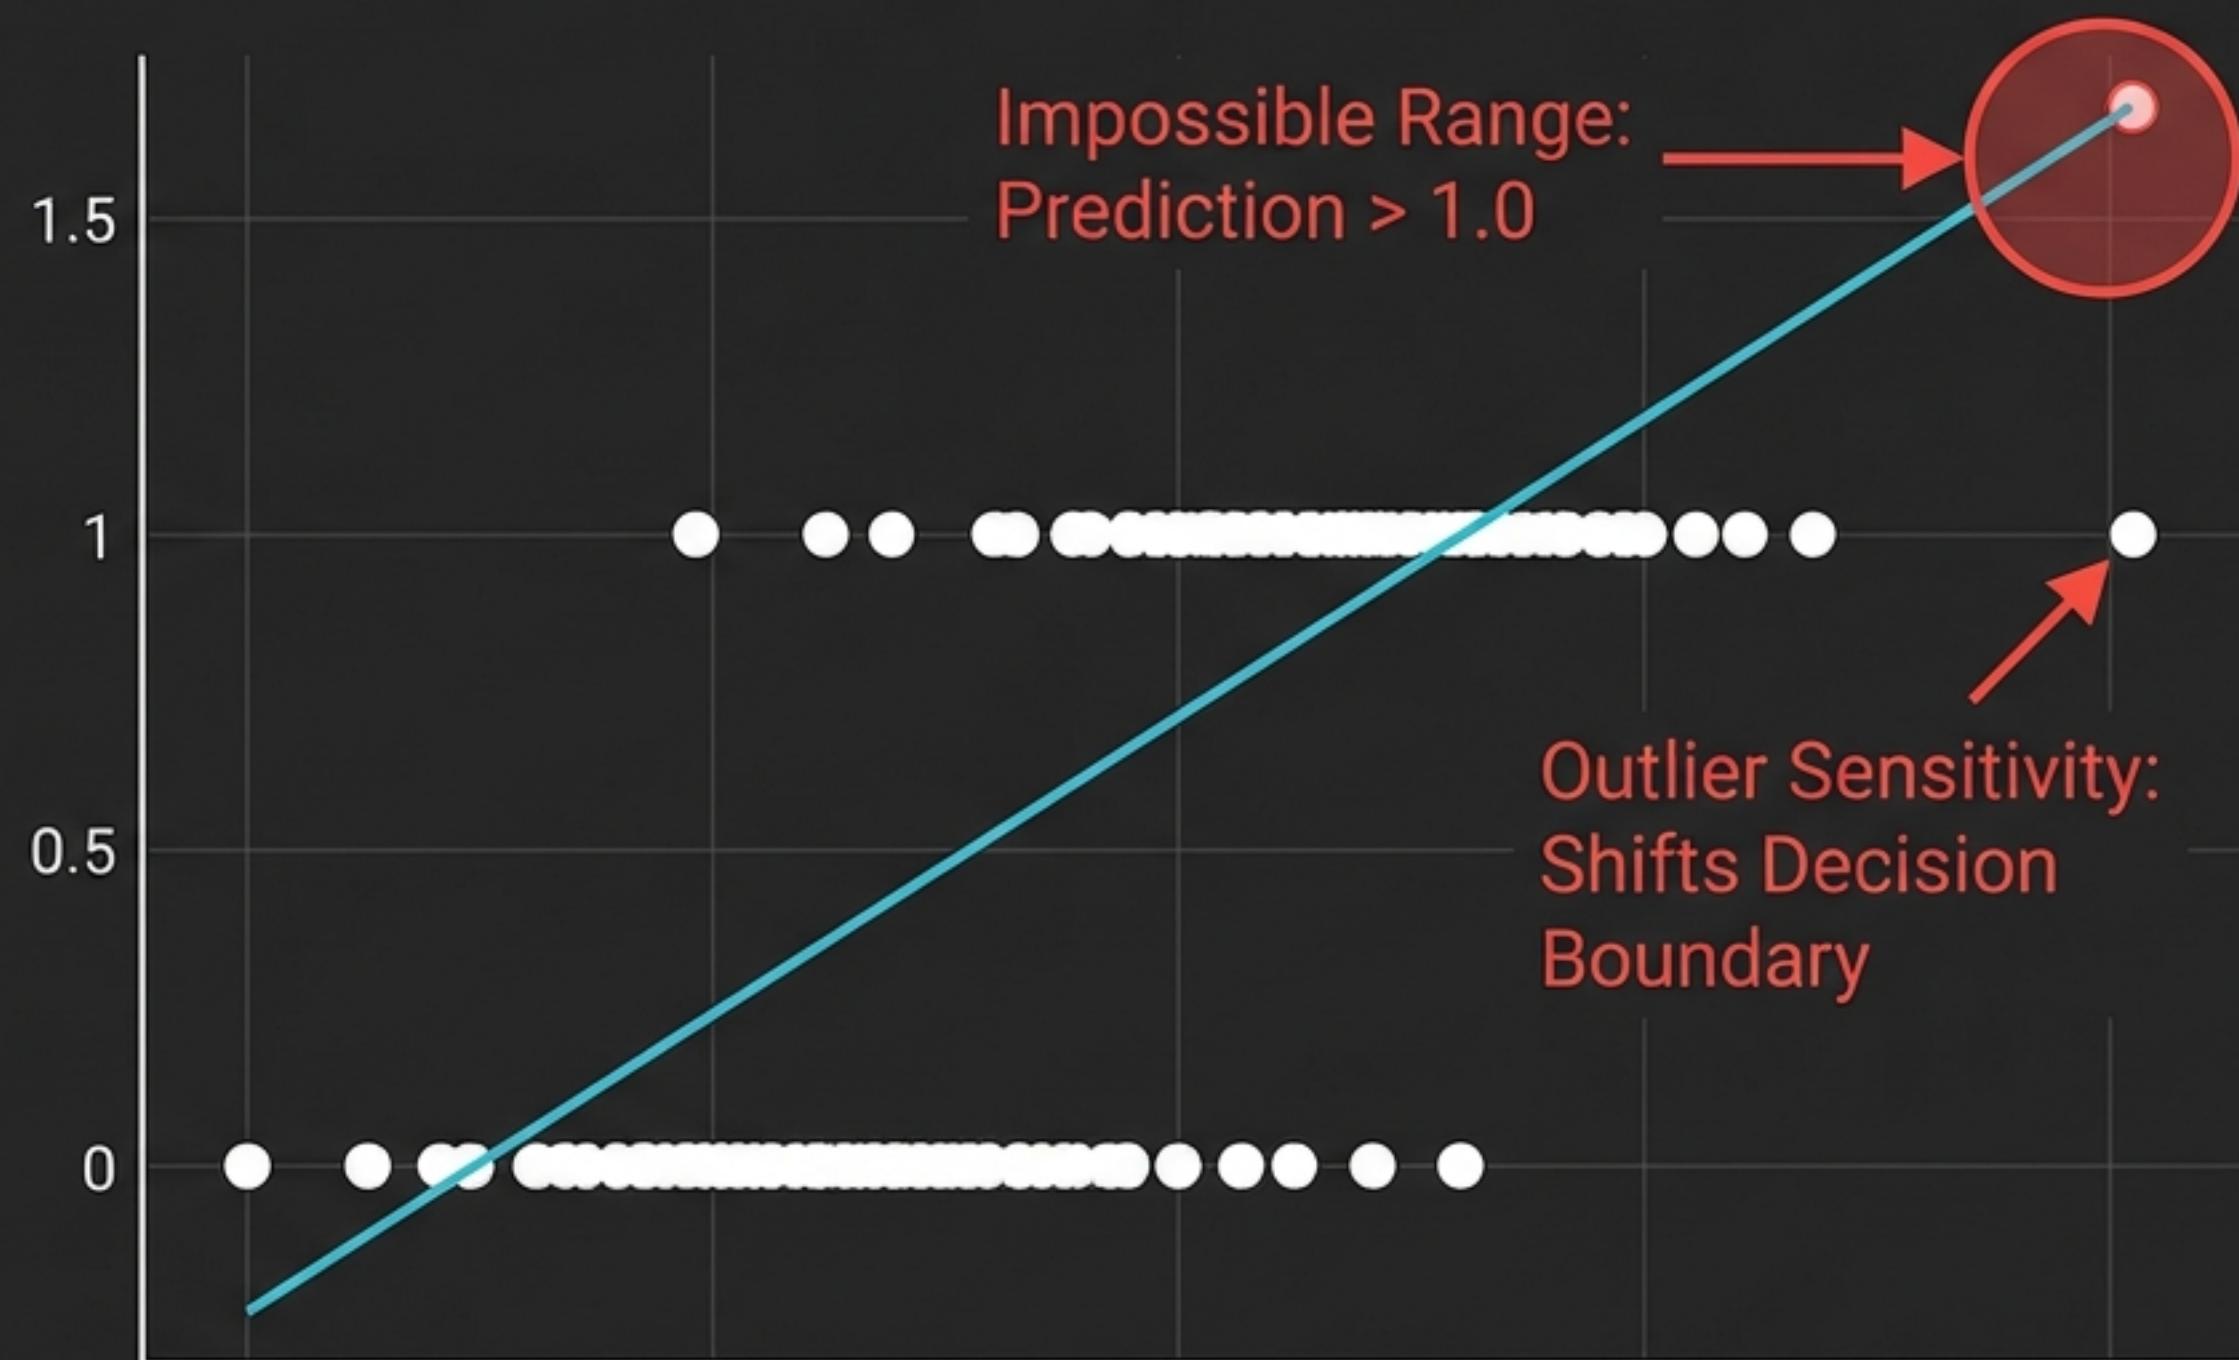
Logistic Regression (Discrete)



The goal shifts from predicting an infinite value to calculating the probability (0 to 1) of a class membership.

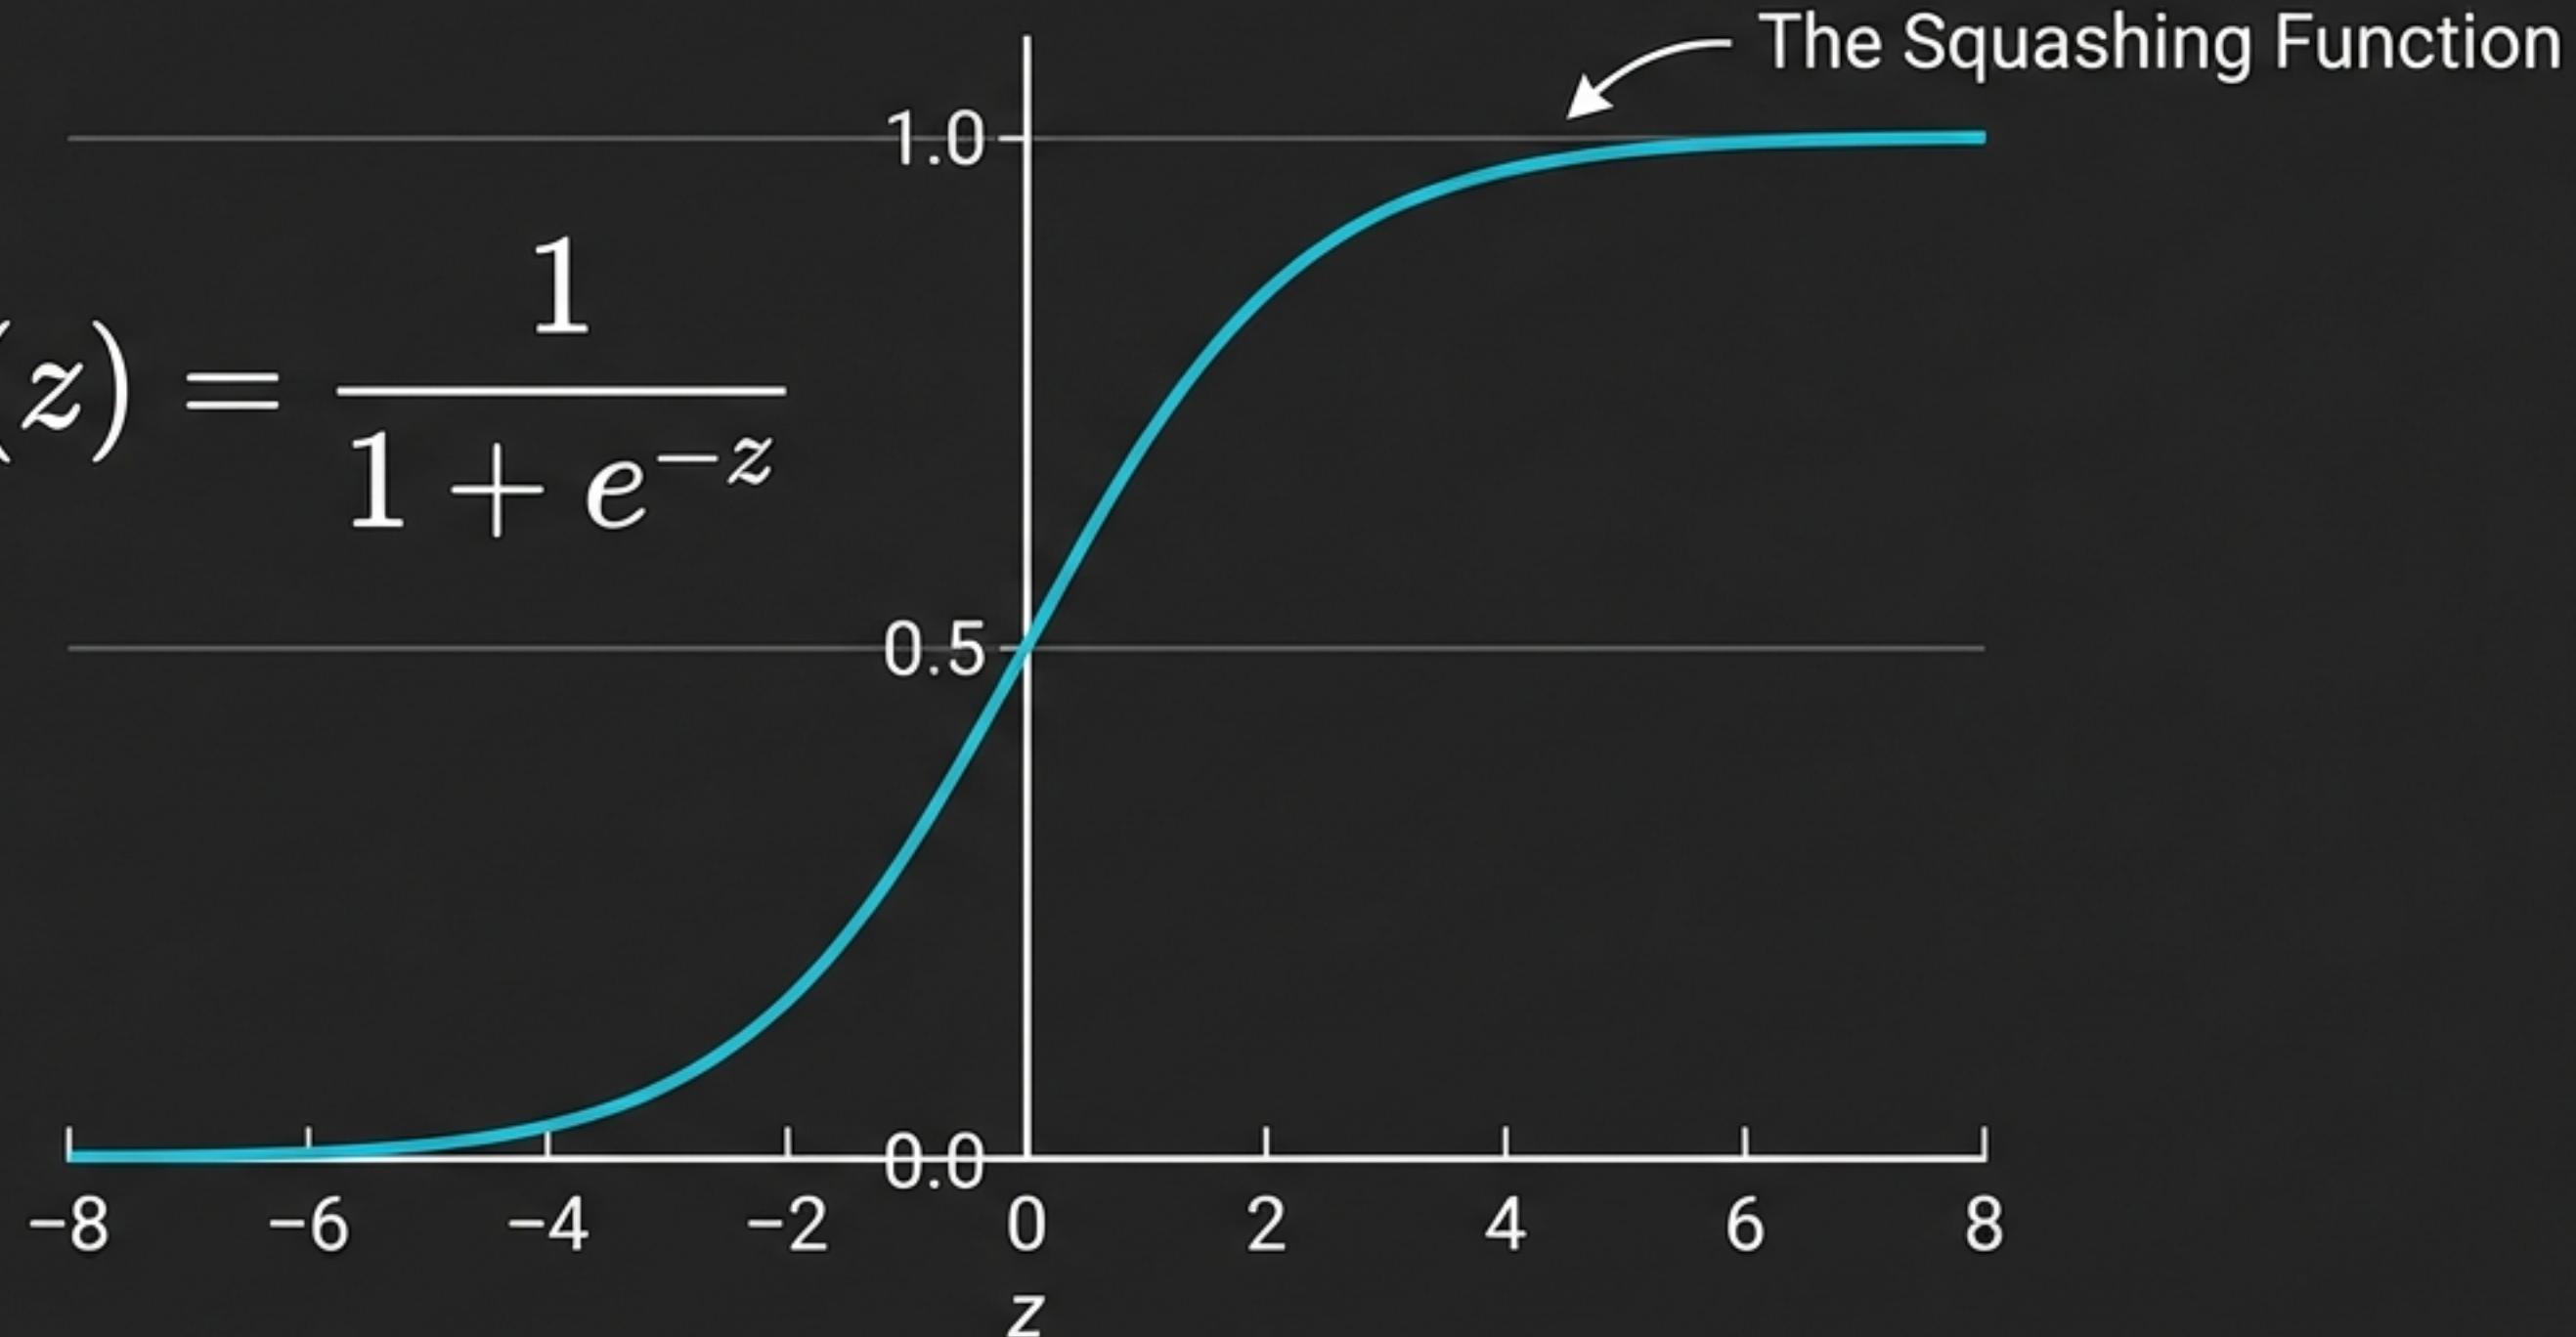
Why Linear Regression Fails at Classification

Linear models are unbounded. A prediction of 1.5 or -0.2 is mathematically invalid for probability.



The Solution: The Sigmoid Function

$$\varphi(z) = \frac{1}{1 + e^{-z}}$$



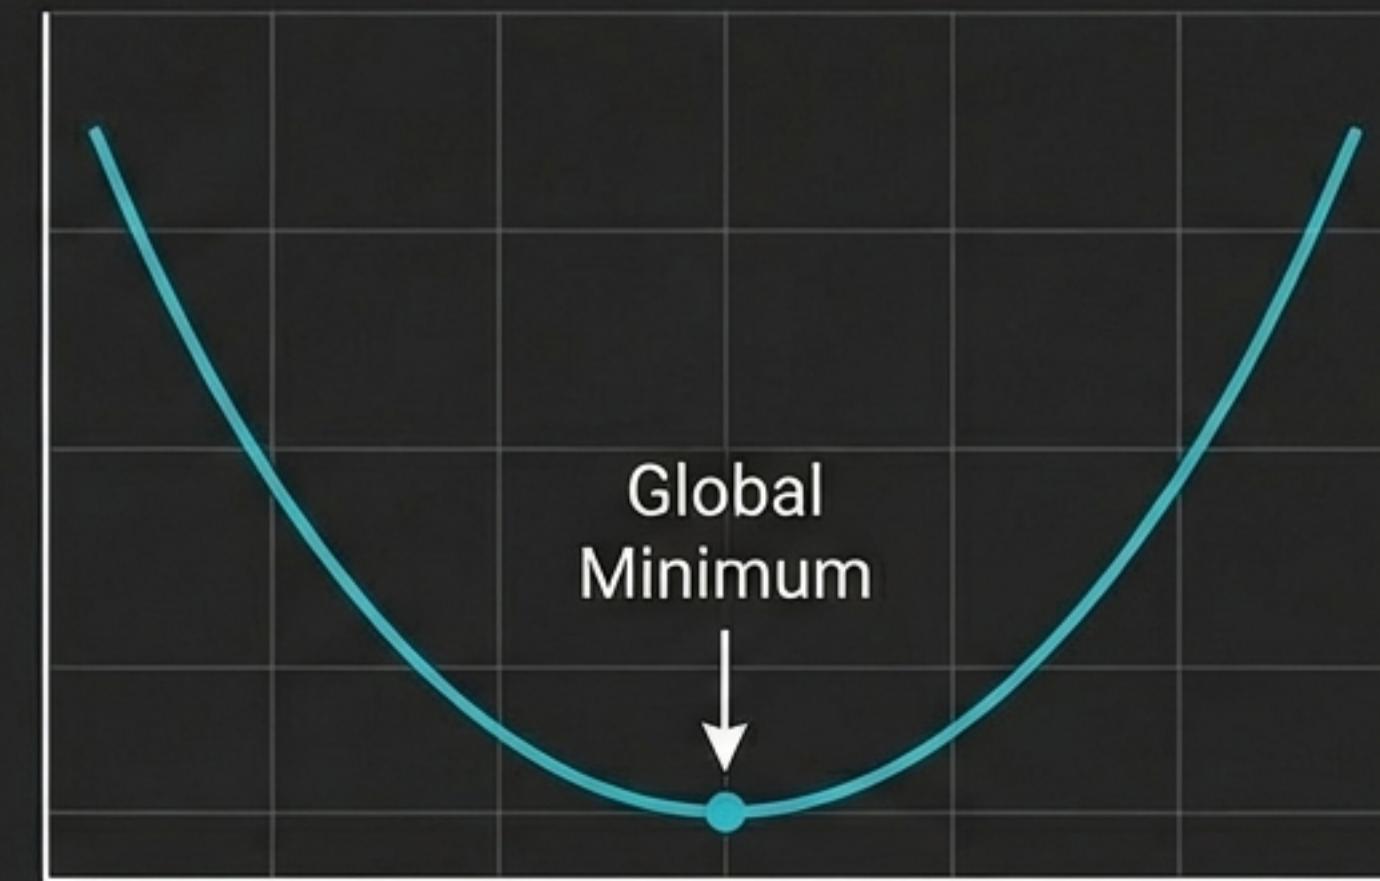
Maps any input z ($-\infty$ to $+\infty$) to a strict probability range of 0 to 1.

The Cost Function: Log Loss

MSE with Logistic



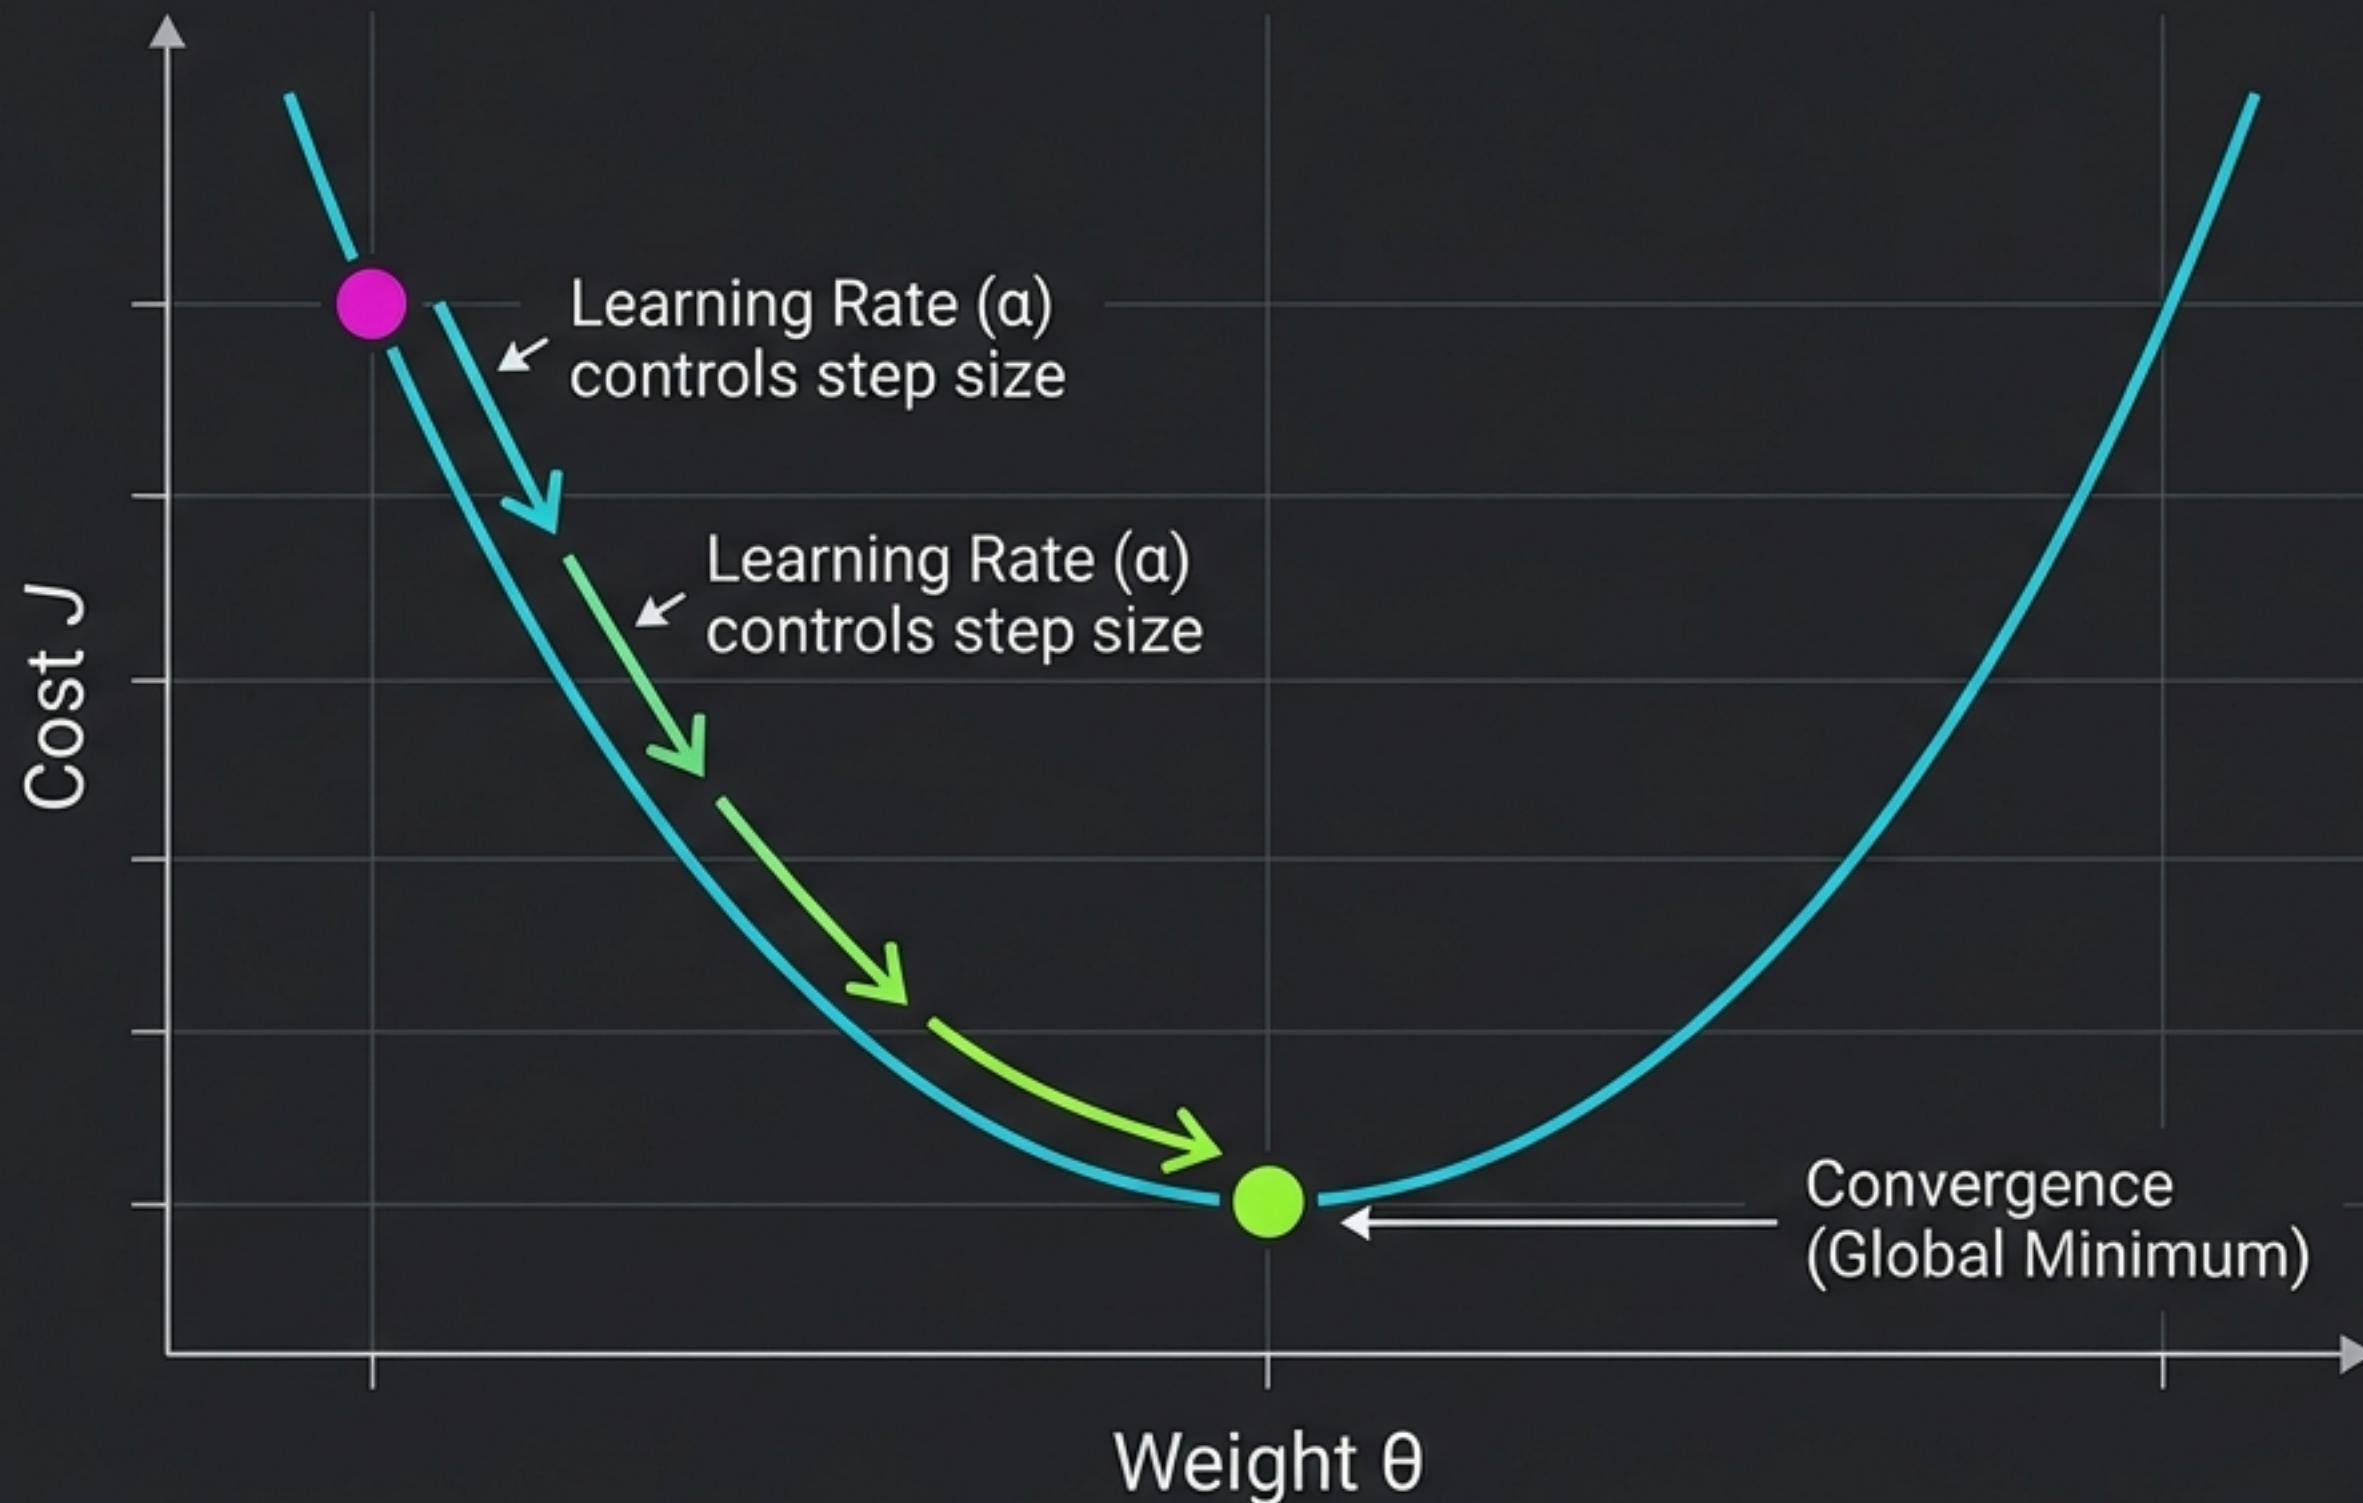
Log Loss / Binary Cross-Entropy



$$\text{Cost} = -y * \log(\hat{y}) - (1-y) * \log(1-\hat{y})$$

We need a convex function (right) to ensure Gradient Descent can always find the lowest error without getting stuck in local valleys.

Reaching Convergence with Gradient Descent



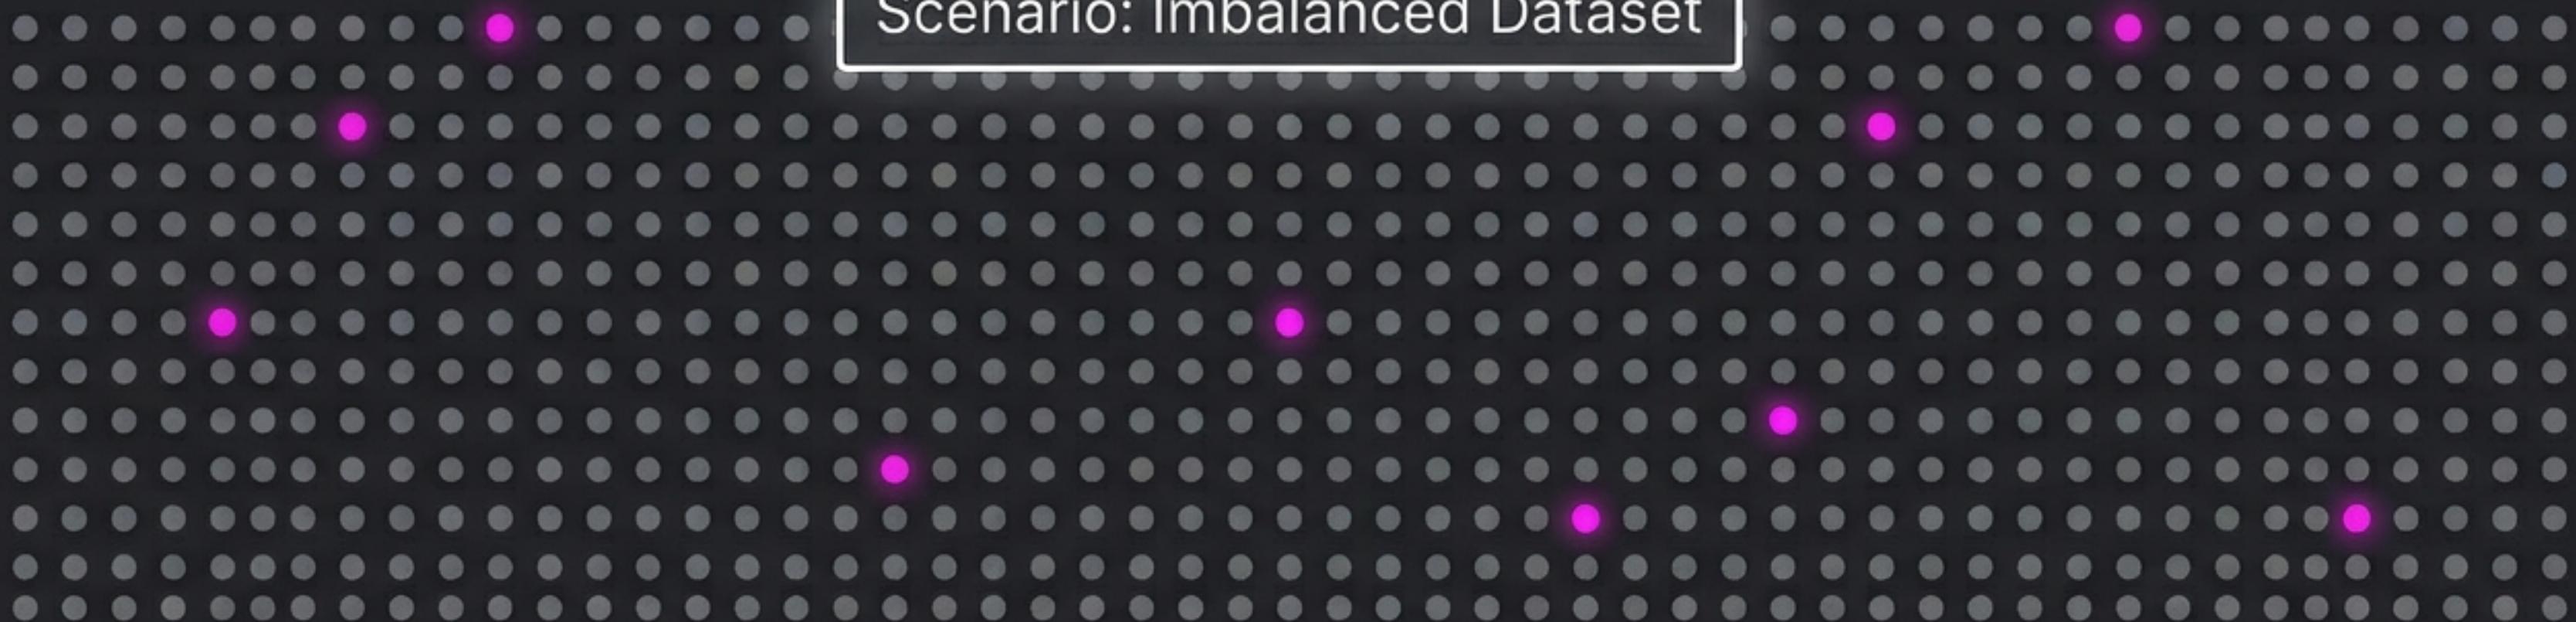
Iteratively adjusts weights (θ) by moving against the gradient of the error.

Beyond Accuracy: The Confusion Matrix

	Actually Positive (1)	Actually Negative (0)
Predicted Positive (1)	 True Positive (TP)	 False Positive (FP) - Type I Error
Predicted Negative (0)	 False Negative (FN) - Type II Error	 True Negative (TN)

The Accuracy Trap

Scenario: Imbalanced Dataset



Model predicts "All Negative"

Captured Positives: 0
Calculated Accuracy: 99%

High Accuracy ≠ Good Model

Precision vs. Recall

Precision Priority
(Spam Filter)

Formula: $TP / (TP + FP)$.

Minimizes False Positives
(False Alarms).

Recall Priority
(Cancer Diagnosis)

Formula: $TP / (TP + FN)$.

Minimizes False Negatives
(Missed Detections).



The F-Beta Score

$$F_{\beta} = \frac{(1 + \beta^2) * (\text{Precision} * \text{Recall})}{(\beta^2 * \text{Precision}) + \text{Recall}}$$

- $\beta = 1$: Harmonic Balance (F1 Score)
- $\beta < 1$: Prioritizes **Precision**
- $\beta > 1$: Prioritizes **Recall**

A single metric to evaluate performance when you need to balance the trade-off.

Implementation: Generating the Data

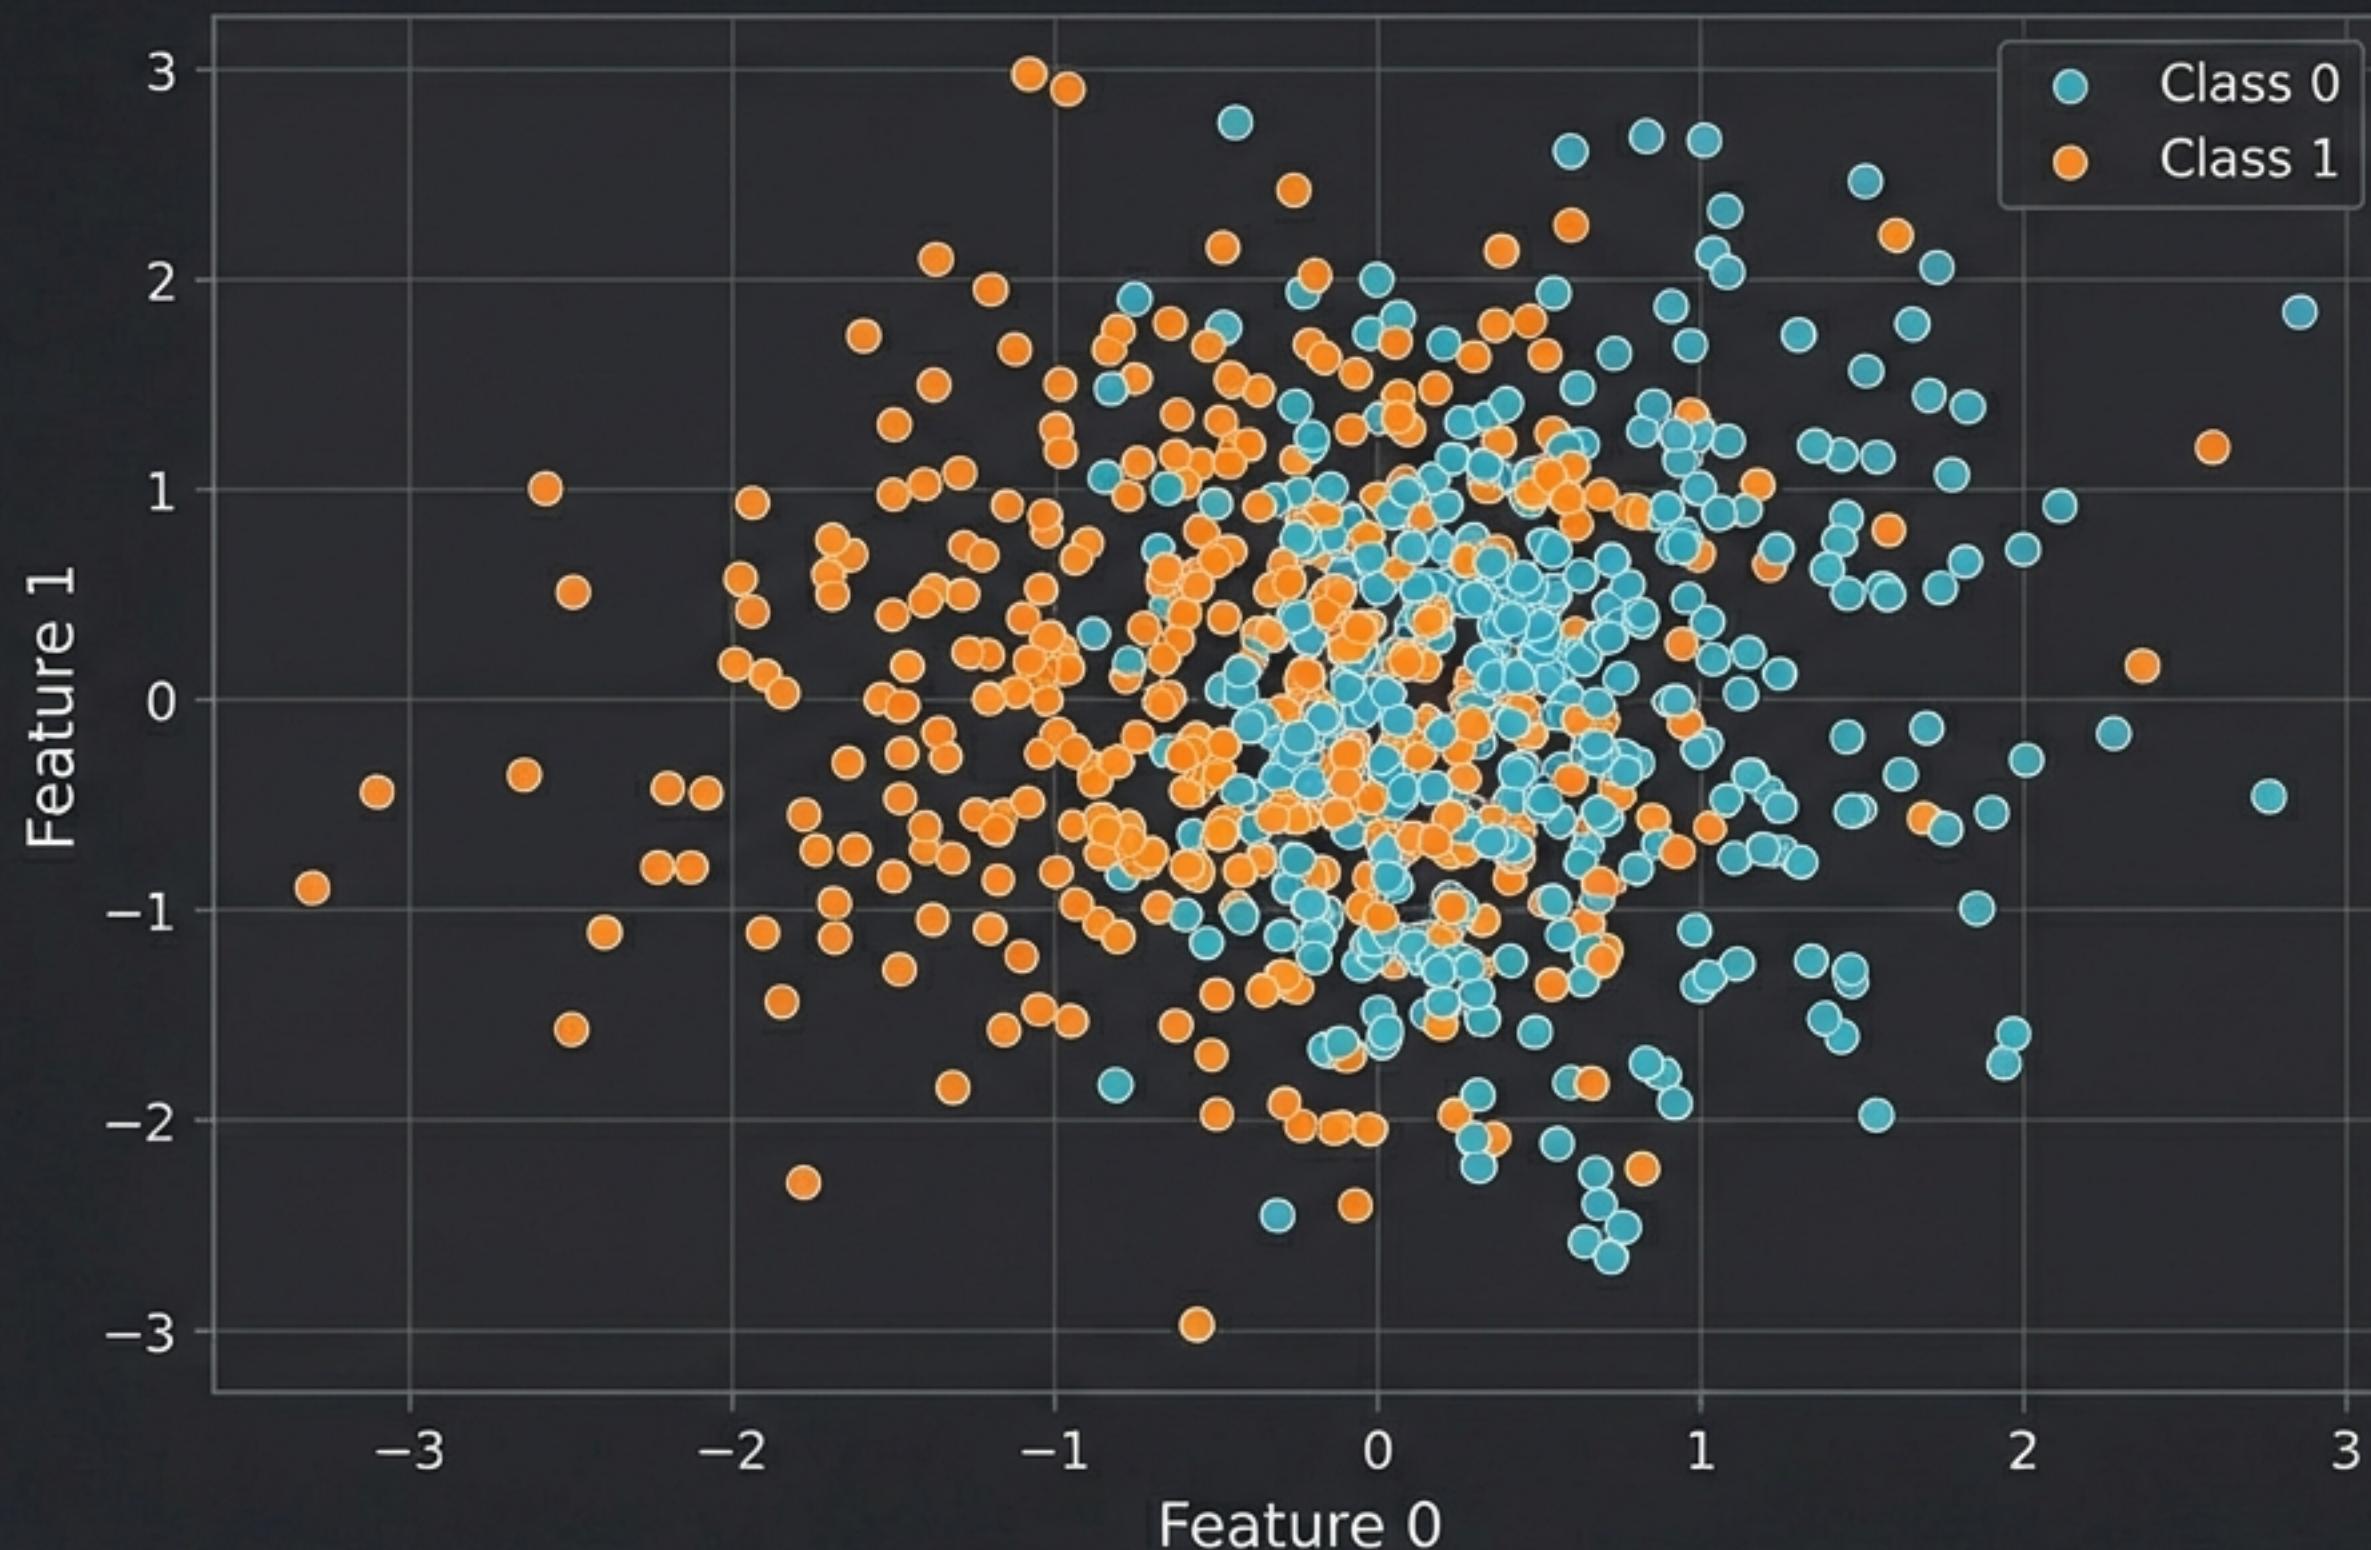
```
train.py  X ▶ ⌂ ...
```

```
1 from sklearn.datasets import make_classification
2 from sklearn.model_selection import train_test_split
3
4 # Generate synthetic binary data
5 X, y = make_classification(n_samples=1000, n_classes=2, n_features=10)
6
7 # Split into training (75%) and testing (25%) sets
8 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
```

Dataset Size

Binary Problem
(0 or 1)

Visualizing the Class Separation

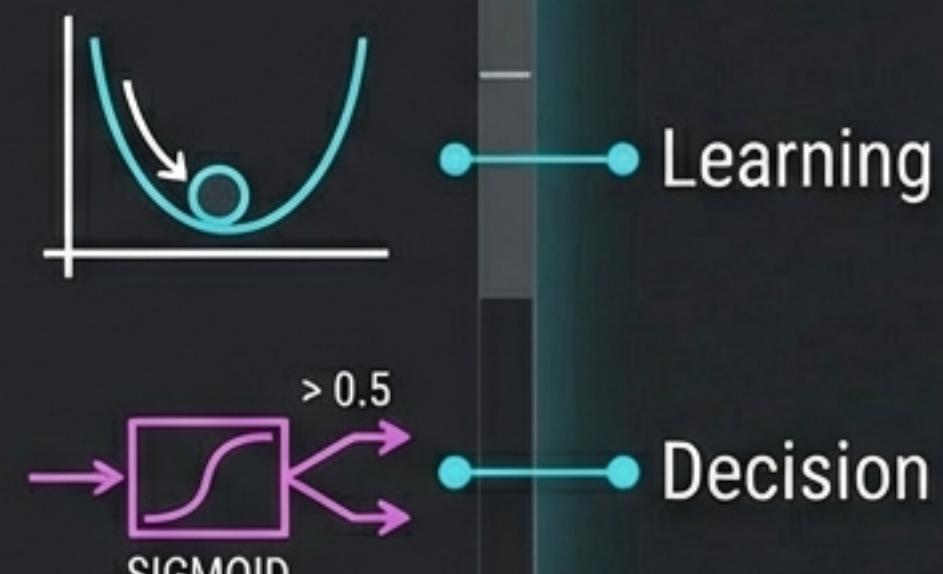


Real-world data is rarely perfectly separable. The model must find the probabilistic boundary.

Training the Model

train.py ×

```
1 from sklearn.linear_model import LogisticRegression  
2  
3 # Initialize  
4 logistic = LogisticRegression()  
5  
6 # 1. Learn the weights (Gradient Descent)  
7 logistic.fit(X_train, y_train)  
8  
9 # 2. Generate predictions (Sigmoid > 0.5)  
10 y_pred = logistic.predict(X_test)
```



Evaluating Performance

```
1      precision    recall   f1-score   support
2
3      0           0.86     0.85      0.86     375
4      1           0.85     0.86      0.86     375
5      1           0.85     0.86      0.86     375
6
7      accuracy                           0.86     750
8
9 Confusion Matrix: [[320, 55], [51, 324]]
```

Result Analysis

Training Accuracy: 0.884

Test Accuracy: 0.858

Stable Model. No significant overfitting detected.

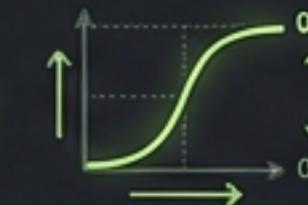
Key Takeaways



Logistic Regression predicts **probability, not continuous values.**



The **Sigmoid Curve** maps input data to a **0-1 range**.



Log Loss (Cost Function) guides the model to the **global minimum**.



Accuracy is deceptive; use **Precision/Recall** for **imbalanced data**.



Scikit-Learn implementation requires just 3 steps: Import, Fit, Predict.

```
...  
import LogisticRegression  
model.fit()  
model.predict()
```